Machine Learning - Homework 2

Pedro Curvo (ist1102716) | Salvador Torpes (ist1102474)

1st Term - 23/24

Pen and Paper Exercises

Dataset

The following dataset will be used for this homework:

D		Input					Output
		y_1	y_2	y_3	y_4	y_5	y_6
Training Observations	x_1	0.24	0.36	1	1	0	A
	x_2	0.16	0.48	1	0	1	A
	x_3	0.32	0.72	0	1	2	A
	x_4	0.54	0.11	0	0	1	B
	x_5	0.66	0.39	0	0	0	B
	x_6	0.76	0.28	1	0	2	B
	x_7	0.41	0.53	0	1	1	B
Testing Observations	x_8	0.38	0.52	0	1	0	A
	x_9	0.42	0.59	0	1	1	B

Table 1: Dataset

1$^{\text{st}}$ Question

a)

In order to build the Bayesian classifier for this dataset, we need to compute the class conditional distributions of $\{y_1, y_2\}$, $\{y_3, y_4\}$ and y_5, which are the groups of independent input variables of our dataset as well as the priors.

Priors   First of all, we will compute the priors $P(y_6 = A)$ and $P(y_6 = B)$:

$$P(y_6 = A) = \frac{3}{7}$$
$$P(y_6 = B) = \frac{4}{7}$$

Distribution of y_1 and y_2

We are told that $y_1 \times y_2 \in \mathbb{R}$ follows a normal 2D distribution. A multivariate normal distribution of m variables $\vec{x} = \{x_1, x_2, ..., x_m\}$ is defined by its mean vector $\vec{\mu}$ and its covariance matrix Σ:

$$P(\vec{x}|\vec{\mu}, \Sigma) = \frac{1}{\sqrt{(2\pi)^m|\Sigma|}} \exp\left(-\frac{1}{2}(\vec{x} - \vec{\mu})^T \cdot \Sigma^{-1} \cdot (\vec{x} - \vec{\mu})\right)$$

In our case, we have $m = 2$, $\vec{x} = \{y_1, y_2\}$ and we need to compute two class conditional distributions $p(\vec{x}|y_6 = A)$ and $p(\vec{x}|y_6 = B)$.

Distribution of $\{y_1, y_2\}$ given $y_6 = A$

Considering the training data in table **??** with class $y_6 = A$, we can compute the mean vector $\vec{\mu}$ and the covariance matrix Σ as follows:

$$\vec{\mu} = \begin{bmatrix} \mu_{y_1} \\ \mu_{y_2} \end{bmatrix} = \frac{1}{3} \cdot \begin{bmatrix} 0.24 + 0.16 + 0.32 \\ 0.36 + 0.48 + 0.72 \end{bmatrix} = \begin{bmatrix} 0.24 \\ 0.52 \end{bmatrix}$$

$$\Sigma = \begin{bmatrix} \sigma_{y_1}^2 & \sigma_{y_1, y_2} \\ \sigma_{y_1, y_2} & \sigma_{y_2}^2 \end{bmatrix} = \frac{1}{3} \cdot \begin{bmatrix} \sum_{i=1}^{3}(y_{1i} - \mu_{y_1})^2 & \sum_{i=1}^{3}(y_{1i} - \mu_{y_1})(y_{2i} - \mu_{y_2}) \\ \sum_{i=1}^{3}(y_{1i} - \mu_{y_1})(y_{2i} - \mu_{y_2}) & \sum_{i=1}^{3}(y_{2i} - \mu_{y_2})^2 \end{bmatrix} = \begin{bmatrix} 0.0043 & 0.0064 \\ 0.0064 & 0.0224 \end{bmatrix}$$

Now we need to compute both $|\Sigma|$ and Σ^{-1}:

$$|\Sigma| = \det \Sigma = 0.0043 \cdot 0.0224 - 0.0064^2 = 5.4613 \cdot 10^{-5}$$

$$\Sigma^{-1} = \begin{bmatrix} 410.156 & -117.188 \\ -117.188 & 78.125 \end{bmatrix}$$

Therefore, we have the normal distribution of $\{y_1, y_2\}$ given $y_6 = A$:

$$P((y_1, y_2)|y_6 = A) = \frac{1}{\sqrt{(2\pi)^2|\Sigma|}} \exp\left(-\frac{1}{2}\left(\begin{bmatrix} y_1 \\ y_2 \end{bmatrix} - \begin{bmatrix} 0.24 \\ 0.52 \end{bmatrix}\right)^T \cdot \begin{bmatrix} 410.156 & -117.188 \\ -117.188 & 78.125 \end{bmatrix} \cdot \left(\begin{bmatrix} y_1 \\ y_2 \end{bmatrix} - \begin{bmatrix} 0.24 \\ 0.52 \end{bmatrix}\right)\right)$$

$$= \frac{1}{\sqrt{(2\pi)^2 \cdot 5.4613 \cdot 10^{-5}}} \exp\left(-\frac{1}{2} \begin{bmatrix} y_1 - 0.24 \\ y_2 - 0.52 \end{bmatrix}^T \cdot \begin{bmatrix} 410.156 & -117.188 \\ -117.188 & 78.125 \end{bmatrix} \cdot \begin{bmatrix} y_1 - 0.24 \\ y_2 - 0.52 \end{bmatrix}\right)$$

Distribution of $\{y_1, y_2\}$ given $y_6 = B$

Considering the training data in table **??** with class $y_6 = B$, we can compute the mean vector $\vec{\mu}$ and the covariance matrix Σ as follows:

$$\vec{\mu} = \begin{bmatrix} \mu_{y_1} \\ \mu_{y_2} \end{bmatrix} = \frac{1}{4} \cdot \begin{bmatrix} 0.54 + 0.66 + 0.76 + 0.41 \\ 0.11 + 0.39 + 0.28 + 0.53 \end{bmatrix} = \begin{bmatrix} 0.5925 \\ 0.3274 \end{bmatrix}$$

$$\Sigma = \begin{bmatrix} \sigma_{y_1}^2 & \sigma_{y_1,y_2} \\ \sigma_{y_1,y_2} & \sigma_{y_2}^2 \end{bmatrix} = \frac{1}{4} \cdot \begin{bmatrix} \sum_{i=1}^{4}(y_{1i} - \mu_{y_1})^2 & \sum_{i=1}^{4}(y_{1i} - \mu_{y_1})(y_{2i} - \mu_{y_2}) \\ \sum_{i=1}^{4}(y_{1i} - \mu_{y_1})(y_{2i} - \mu_{y_2}) & \sum_{i=1}^{4}(y_{2i} - \mu_{y_2})^2 \end{bmatrix} = \begin{bmatrix} 0.0171 & -0.0073 \\ -0.0073 & 0.0236 \end{bmatrix}$$

Now we need to compute both $|\Sigma|$ and Σ^{-1}:

$$|\Sigma| = \det \Sigma = 0.0075 \cdot 0.0075 - (-0.0025)^2 = 3.519 \cdot 10^{-4}$$

$$\Sigma^{-1} = \begin{bmatrix} 67.1101 & 20.7954 \\ 20.7954 & 48.7831 \end{bmatrix}$$

Therefore, we have the normal distribution of $\{y_1, y_2\}$ given $y_6 = B$:

$$P((y_1, y_2)|y_6 = B) = \frac{1}{\sqrt{(2\pi)^2 |\Sigma|}} \exp\left(-\frac{1}{2}\left(\begin{bmatrix} y_1 \\ y_2 \end{bmatrix} - \begin{bmatrix} 0.5925 \\ 0.3274 \end{bmatrix}\right)^T \cdot \begin{bmatrix} 67.1101 & 20.7954 \\ 20.7954 & 48.7831 \end{bmatrix} \cdot \left(\begin{bmatrix} y_1 \\ y_2 \end{bmatrix} - \begin{bmatrix} 0.5925 \\ 0.3274 \end{bmatrix}\right)\right)$$

$$= \frac{1}{\sqrt{(2\pi)^2 \cdot 3.519 \cdot 10^{-4}}} \exp\left(-\frac{1}{2} \begin{bmatrix} y_1 - 0.5925 \\ y_2 - 0.3274 \end{bmatrix}^T \cdot \begin{bmatrix} 67.1101 & 20.7954 \\ 20.7954 & 48.7831 \end{bmatrix} \cdot \begin{bmatrix} y_1 - 0.5925 \\ y_2 - 0.3274 \end{bmatrix}\right)$$

Distribution of y_3 and y_4

The class conditional distributions of y_3 and y_4 come directly from the information in table **??** and they are given by:

$P(y_3 \cap y_4 | y_6 = A)$		y_3	
		0	1
y_4	0	$P(y_3 = 0 \cap y_4 = 0 | y_6 = A) = 0$	$P(y_3 = 1 \cap y_4 = 0 | y_6 = A) = \frac{1}{3}$
	1	$P(y_3 = 0 \cap y_4 = 1 | y_6 = A) = \frac{1}{3}$	$P(y_3 = 1 \cap y_4 = 1 | y_6 = A) = \frac{1}{3}$

Table 2: Distribution of y_3 and y_4 given $y_6 = A$

$P(y_3 \cap y_4 \| y_6 = B)$		y_3	
		0	1
y_4	0	$P(y_3 = 0 \cap y_4 = 0 \| y_6 = B) = \frac{1}{2}$	$P(y_3 = 1 \cap y_4 = 0 \| y_6 = B) = \frac{1}{4}$
	1	$P(y_3 = 0 \cap y_4 = 1 \| y_6 = B) = \frac{1}{4}$	$P(y_3 = 1 \cap y_4 = 1 \| y_6 = B) = 0$

Table 3: Distribution of y_3 and y_4 given $y_6 = B$

Distribution of y_5

The class conditional distribution of y_5 is given by:

$P(y_5 \| y_6)$		y_5		
		0	1	2
y_6	A	$P(y_5 = 0 \| y_6 = A) = \frac{1}{3}$	$P(y_5 = 1 \| y_6 = A) = \frac{1}{3}$	$P(y_5 = 2 \| y_6 = A) = \frac{1}{3}$
	B	$P(y_5 = 0 \| y_6 = B) = \frac{1}{4}$	$P(y_5 = 1 \| y_6 = B) = \frac{1}{2}$	$P(y_5 = 2 \| y_6 = B) = \frac{1}{4}$

Table 4: Distribution of y_5 given y_6

b)

In order to classify the testing observations, we will need to compute the posterior probabilities. Under a MAP assumption, the predicted class for each testing observation is the one that maximizes the posterior probability. Since we are only interested in the maximum value over all classes, we can ignore the denominator of the posterior probability formula. We have two testing observations, x_8 and x_9, and we will compute the posterior probabilities for each of them:

Posterior probabilities for x_8

This training observation has the following values for the input variables: $y_1 = 0.38$, $y_2 = 0.52$, $y_3 = 0$, $y_4 = 1$ and $y_5 = 0$.

$$P(y_6 = A | x_8) = \frac{P(x_8 | y_6 = A) \cdot P(y_6 = A)}{P(x_8)} \propto P(x_8 | y_6 = A) \cdot P(y_6 = A) =$$

$$= P(y_1 = 0.38, y_2 = 0.52 | y_6 = A) \cdot P(y_3 = 0, y_4 = 1 | y_6 = A) \cdot P(y_5 = 0 | y_6 = A) \cdot P(y_6 = A) =$$

$$= \frac{1}{\sqrt{(2\pi)^2 \cdot 5.4613 \cdot 10^{-5}}} \exp\left(-\frac{1}{2} \begin{bmatrix} 0.38 - 0.24 \\ 0.52 - 0.52 \end{bmatrix}^T \cdot \begin{bmatrix} 410.156 & -117.188 \\ -117.188 & 78.125 \end{bmatrix} \cdot \begin{bmatrix} 0.38 - 0.24 \\ 0.52 - 0.52 \end{bmatrix}\right) \cdot \frac{1}{3} \cdot \frac{1}{3} \cdot \frac{3}{7} =$$

$$= 0.3868 \cdot \frac{1}{3} \cdot \frac{1}{3} \cdot \frac{3}{7} = 0.01842$$

$$P(y_6 = B|x_8) = \frac{P(x_8|y_6 = B) \cdot P(y_6 = B)}{P(x_8)} \propto P(x_8|y_6 = B) \cdot P(y_6 = B) =$$

$$= P(y_1 = 0.38, y_2 = 0.52|y_6 = B) \cdot P(y_3 = 0, y_4 = 1|y_6 = B) \cdot P(y_5 = 0|y_6 = B) \cdot P(y_6 = B) =$$

$$= \frac{1}{\sqrt{(2\pi)^2 \cdot 3.519 \cdot 10^{-4}}} \exp\left(-\frac{1}{2}\begin{bmatrix} 0.38 - 0.5925 \\ 0.52 - 0.3274 \end{bmatrix}^T \cdot \begin{bmatrix} 67.1101 & 20.7954 \\ 20.7954 & 48.7831 \end{bmatrix} \cdot \begin{bmatrix} 0.38 - 0.5925 \\ 0.52 - 0.3274 \end{bmatrix}\right) \cdot \frac{1}{4} \cdot \frac{1}{4} \cdot \frac{4}{7} =$$

$$= 1.7677 \cdot \frac{1}{4} \cdot \frac{1}{4} \cdot \frac{4}{7} = 0.06313$$

Posterior probabilities for x_9

This training observation has the following values for the input variables: $y_1 = 0.42$, $y_2 = 0.59$, $y_3 = 0$, $y_4 = 1$ and $y_5 = 1$.

$$P(y_6 = A|x_9) = \frac{P(x_9|y_6 = A) \cdot P(y_6 = A)}{P(x_9)} \propto P(x_9|y_6 = A) \cdot P(y_6 = A) =$$

$$= P(y_1 = 0.42, y_2 = 0.59|y_6 = A) \cdot P(y_3 = 0, y_4 = 1|y_6 = A) \cdot P(y_5 = 1|y_6 = A) \cdot P(y_6 = A) =$$

$$= \frac{1}{\sqrt{(2\pi)^2 \cdot 5.4613 \cdot 10^{-5}}} \exp\left(-\frac{1}{2}\begin{bmatrix} 0.42 - 0.24 \\ 0.59 - 0.52 \end{bmatrix}^T \cdot \begin{bmatrix} 410.156 & -117.188 \\ -117.188 & 78.125 \end{bmatrix} \cdot \begin{bmatrix} 0.42 - 0.24 \\ 0.59 - 0.52 \end{bmatrix}\right) \cdot \frac{1}{3} \cdot \frac{1}{3} \cdot \frac{3}{7} =$$

$$= 0.1013 \cdot \frac{1}{3} \cdot \frac{1}{3} \cdot \frac{3}{7} = 0.00482$$

$$P(y_6 = B|x_9) = \frac{P(x_9|y_6 = B) \cdot P(y_6 = B)}{P(x_9)} \propto P(x_9|y_6 = B) \cdot P(y_6 = B) =$$

$$= P(y_1 = 0.42, y_2 = 0.59|y_6 = B) \cdot P(y_3 = 0, y_4 = 1|y_6 = B) \cdot P(y_5 = 1|y_6 = B) \cdot P(y_6 = B) =$$

$$= \frac{1}{\sqrt{(2\pi)^2 \cdot 3.519 \cdot 10^{-4}}} \exp\left(-\frac{1}{2}\begin{bmatrix} 0.42 - 0.5925 \\ 0.59 - 0.3274 \end{bmatrix}^T \cdot \begin{bmatrix} 67.1101 & 20.7954 \\ 20.7954 & 48.7831 \end{bmatrix} \cdot \begin{bmatrix} 0.42 - 0.5925 \\ 0.59 - 0.3274 \end{bmatrix}\right) \cdot \frac{1}{4} \cdot \frac{1}{2} \cdot \frac{4}{7} =$$

$$= 1.4927 \cdot \frac{1}{4} \cdot \frac{1}{2} \cdot \frac{4}{7} = 0.05331$$

Predicted classes

Organizing the posterior probabilities in a table, we have:

Observation	$P(y_6 = A|x_i)$	$P(y_6 = B|x_i)$
x_8	0.01842	**0.06313**
x_9	0.00482	**0.05331**

Table 5: Posterior probabilities for the testing observations

Therefore, the predicted class for both x_8 and x_9 is $y_6 = B$.

c)

Let's consider the following classifier with a unknown treshold θ whose value we aim to find:

$$f(x|\theta) = \begin{cases} A & \text{if } P(y_6 = A|x) > \theta \\ B & \text{if } P(y_6 = A|x) \leq \theta \end{cases}$$

Finding $P(y_6 = A|x)$

We are now working under a ML assumption, so, in order to classify each test observation we will only need the conditional distribution $P(x|y_6 = A)$ because every class has the same prior probability. We used the values in the previous section and divided them by the corresponding a priori probability:

$$P(y_6 = A|x_8) = 0.04298$$

$$P(y_6 = A|x_9) = 0.01126$$

Accuracy The accuracy of a classifier is given by:

$$\text{Accuracy} = \frac{TP + TN}{TP + TN + FP + FN}$$

where TP is the number of true positives, TN is the number of true negatives, FP is the number of false positives and FN is the number of false negatives. In our case, accuracy can only have three possible values: 0, 0.5 or 1: 0 when both x_8 and x_9 are misclassified, 0.5 when only one of them is misclassified and 1 when both are correctly classified. In order to maximize the accuracy, we want it to be 1. We know from table **??** that x_8 belongs to class A and x_9 belongs to class B. Therefore, in order to maximize the accuracy, our classifier needs to classify x_8 as A and x_9 as B. $f(x_8|\theta) = A$ therefore $\theta < P(y_6 = A|x_8) = 0.04298$ and $f(x_9|\theta) = B$ therefore $\theta > P(y_6 = A|x_9) = 0.01126$.

$$\theta \in]0.01126; 0.04298[$$

Any value of θ in this interval will maximize the accuracy of our classifier.

2$^{\text{nd}}$ Question

a)

In order to obtain y_2 under an equal-range discretization, we followed the rule:

$$y_{2_{\text{normalized}}} = \begin{cases} 0 & \text{if } y_2 \in [0, 0.5) \\ 1 & \text{if } y_2 \in [0.5, 1] \end{cases}$$

The normalized values are:

Dataset	x_1	x_2	x_3	x_4	x_5	x_6	x_7	x_8	x_9
y_2	0	0	1	0	0	0	1	1	1

Table 6: Normalized y_2 values

y_1 will now be considered an output variable and y_2 to y_6 will be considered input variables. Considering the normalized values of y_2, we can rewrite the dataset as follows:

D		Output	Input				
Fold		y_1	$y_{2_{norm}}$	y_3	y_4	y_5	y_6
F_1	x_1	0.24	0	1	1	0	A
	x_2	0.16	0	1	0	1	A
	x_3	0.32	1	0	1	2	A
F_2	x_4	0.54	0	0	0	1	B
	x_5	0.66	0	0	0	0	B
	x_6	0.76	0	1	0	2	B
F_3	x_7	0.41	1	0	1	1	B
	x_8	0.38	1	0	1	0	A
	x_9	0.42	1	0	1	1	B

Table 7: Dataset D divided into three folds

Adittionally, we have divided the dataset into three folds, F_1, F_2 and F_3.

b)

In this exercise we aim to compute a kNN (k nearest neighbors - Lazy Learning) classifier considering the following parameters:

- $k = 3$

- **Hamming Distance** as the distance to be used to compute the nearest neighbors of a given observation.

$$d_H(x_i, x_j) = \sum_{l=1}^{m} \delta(y_{il}, y_{jl})$$

where m is the number of input variables and y_{il} is the value of the l^{th} input variable of the i^{th} observation.

We have divided our dataset in folds in order to perform a cross validation. We will only be interested in the first iteration of the cross validation, where F_3 is the testing fold and F_1 and F_2 are the training folds. We will now compute the kNN classifier for each observation in F_3 and afterwards we will compute the MAE (Mean Absolute Error) for the testing fold.

Computing the Hamming Distances

Testing Observation (x_i)	$d_H(x_i, x_1)$	$d_H(x_i, x_2)$	$d_H(x_i, x_3)$	$d_H(x_i, x_4)$	$d_H(x_i, x_5)$	$d_H(x_i, x_6)$
x_7	4	4	2	2	3	4
x_8	2	4	1	4	3	5
x_9	4	4	2	2	3	4

Table 8: Hamming distances between the testing observation x_i and the training observations x_j

We are considering $k = 3$, so we will only need the three nearest neighbors of each testing observation. In the table above we have filled with yellow the three nearest neighbors of each testing observation.

Predicted value of y_1 for each testing observation

The output value we are working with is numerical, so the predicted value of y_1 for each testing observation will be:

$$\hat{y}_{1j} = \frac{\sum_{i=1}^{k} \frac{1}{d_H(x_i, x_j)} \cdot y_{1j}}{\sum_{i=1}^{k} \frac{1}{d_H(x_i, x_j)}}$$

where k is the number of nearest neighbors, $d_H(x_i, x_j)$ is the Hamming distance between the testing observation x_i and the i^{th} nearest neighbor and y_{1j} is the value of the output variable of the j^{th} nearest neighbor. $\frac{1}{d_H(x_i, x_j)}$ is the weight of the i^{th} nearest neighbor.

$$\hat{y}_{17} = \frac{1}{\frac{1}{2} + \frac{1}{2} + \frac{1}{3}} \cdot \left(\frac{1}{2} \cdot 0.32 + \frac{1}{2} \cdot 0.54 + \frac{1}{3} \cdot 0.66 \right) = 0.4875$$

$$\hat{y}_{18} = \frac{1}{\frac{1}{2} + \frac{1}{1} + \frac{1}{3}} \cdot \left(\frac{1}{2} \cdot 0.24 + 1 \cdot 0.32 + \frac{1}{3} \cdot 0.66 \right) = 0.36$$

$$\hat{y}_{19} = \frac{1}{\frac{1}{2} + \frac{1}{2} + \frac{1}{3}} \cdot \left(\frac{1}{2} \cdot 0.32 + \frac{1}{2} \cdot 0.54 + \frac{1}{3} \cdot 0.66 \right) = 0.4875$$

We have the following predicted values for y_1:

Testing Observation (x_i)	$\hat{y}_{1i}$	y_{1i}
x_7	0.4875	0.41
x_8	0.36	0.38
x_9	0.4875	0.42

Table 9: Predicted values of y_1 for each testing observation

MAE

The MAE is given by:

$$\text{MAE} = \frac{1}{n} \sum_{i=1}^{n} |y_{1i} - \hat{y_{1i}}|$$

where n is the number of testing observations, y_{1i} is the value of the output variable of the i^{th} testing observation and $\hat{y_{1i}}$ is the predicted value of the output variable of the i^{th} testing observation.

$$\text{MAE} = \frac{1}{3} \cdot (|0.41 - 0.4875| + |0.38 - 0.36| + |0.42 - 0.4875|) = 0.055$$

Programming and Critical Analysis

1ˢᵗ Question

2ⁿᵈ Question

3ʳᵈ Question

aa