Machine Learning - Homework 2

Pedro Curvo (ist1102716) | Salvador Torpes (ist1102474)

1st Term - 23/24

Pen and Paper Exercises

Dataset

The following dataset will be used for this homework:

D		Input					Output
		y_1	y_2	y_3	y_4	y_5	y_6
Training Observations	x_1	0.24	0.36	1	1	0	A
	x_2	0.16	0.48	1	0	1	A
	x_3	0.32	0.72	0	1	2	A
	x_4	0.54	0.11	0	0	1	B
	x_5	0.66	0.39	0	0	0	B
	x_6	0.76	0.28	1	0	2	B
	x_7	0.41	0.53	0	1	1	B
Testing Observations	x_8	0.38	0.52	0	1	0	A
	x_9	0.42	0.59	0	1	1	B

Table 1: Dataset

1$^{\text{st}}$ Question

a)

In order to build the Bayesian classifier for this dataset, we need to compute the class conditional distributions of $\{y_1, y_2\}$, $\{y_3, y_4\}$ and y_5, which are the groups of independent input variables of our dataset as well as the priors.

Priors   First of all, we will compute the priors $P(y_6 = A)$ and $P(y_6 = B)$:

$$P(y_6 = A) = \frac{3}{7}$$
$$P(y_6 = B) = \frac{4}{7}$$

Distribution of y_1 and y_2

We are told that $y_1 \times y_2 \in \mathbb{R}$ follows a normal 2D distribution. A multivariate normal distribution of m variables $\vec{x} = \{x_1, x_2, ..., x_m\}$ is defined by its mean vector $\vec{\mu}$ and its covariance matrix Σ:

$$P(\vec{x}|\vec{\mu}, \Sigma) = \frac{1}{\sqrt{(2\pi)^m |\Sigma|}} \exp\left(-\frac{1}{2}(\vec{x} - \vec{\mu})^T \cdot \Sigma^{-1} \cdot (\vec{x} - \vec{\mu}) \right)$$

In our case, we have $m = 2$, $\vec{x} = \{y_1, y_2\}$ and we need to compute two class conditional distributions $p(\vec{x}|y_6 = A)$ and $p(\vec{x}|y_6 = B)$.

Distribution of $\{y_1, y_2\}$ given $y_6 = A$

Considering the training data in table **??** with class $y_6 = A$, we can compute the mean vector $\vec{\mu}$ and the covariance matrix Σ as follows:

$$\vec{\mu} = \begin{bmatrix} \mu_{y_1} \\ \mu_{y_2} \end{bmatrix} = \frac{1}{3} \cdot \begin{bmatrix} 0.24 + 0.16 + 0.32 \\ 0.36 + 0.48 + 0.72 \end{bmatrix} = \begin{bmatrix} 0.24 \\ 0.52 \end{bmatrix}$$

$$\Sigma = \begin{bmatrix} \sigma_{y_1}^2 & \sigma_{y_1, y_2} \\ \sigma_{y_1, y_2} & \sigma_{y_2}^2 \end{bmatrix} = \frac{1}{3} \cdot \begin{bmatrix} \sum_{i=1}^{3}(y_{1i} - \mu_{y_1})^2 & \sum_{i=1}^{3}(y_{1i} - \mu_{y_1})(y_{2i} - \mu_{y_2}) \\ \sum_{i=1}^{3}(y_{1i} - \mu_{y_1})(y_{2i} - \mu_{y_2}) & \sum_{i=1}^{3}(y_{2i} - \mu_{y_2})^2 \end{bmatrix} = \begin{bmatrix} 0.0043 & 0.0064 \\ 0.0064 & 0.0224 \end{bmatrix}$$

Now we need to compute both $|\Sigma|$ and Σ^{-1}:

$$|\Sigma| = \det \Sigma = 0.0043 \cdot 0.0224 - 0.0064^2 = 5.4613 \cdot 10^{-5}$$

$$\Sigma^{-1} = \begin{bmatrix} 410.156 & -117.188 \\ -117.188 & 78.125 \end{bmatrix}$$

Therefore, we have the normal distribution of $\{y_1, y_2\}$ given $y_6 = A$:

$$P((y_1, y_2)|y_6 = A) = \frac{1}{\sqrt{(2\pi)^2 |\Sigma|}} \exp\left(-\frac{1}{2}\left(\begin{bmatrix} y_1 \\ y_2 \end{bmatrix} - \begin{bmatrix} 0.24 \\ 0.52 \end{bmatrix} \right)^T \cdot \begin{bmatrix} 410.156 & -117.188 \\ -117.188 & 78.125 \end{bmatrix} \cdot \left(\begin{bmatrix} y_1 \\ y_2 \end{bmatrix} - \begin{bmatrix} 0.24 \\ 0.52 \end{bmatrix} \right) \right)$$

$$= \frac{1}{\sqrt{(2\pi)^2 \cdot 5.4613 \cdot 10^{-5}}} \exp\left(-\frac{1}{2} \begin{bmatrix} y_1 - 0.24 \\ y_2 - 0.52 \end{bmatrix}^T \cdot \begin{bmatrix} 410.156 & -117.188 \\ -117.188 & 78.125 \end{bmatrix} \cdot \begin{bmatrix} y_1 - 0.24 \\ y_2 - 0.52 \end{bmatrix}\right)$$

Distribution of $\{y_1, y_2\}$ given $y_6 = B$

Considering the training data in table **??** with class $y_6 = B$, we can compute the mean vector $\vec{\mu}$ and the covariance matrix Σ as follows:

$$\vec{\mu} = \begin{bmatrix} \mu_{y_1} \\ \mu_{y_2} \end{bmatrix} = \frac{1}{4} \cdot \begin{bmatrix} 0.54 + 0.66 + 0.76 + 0.41 \\ 0.11 + 0.39 + 0.28 + 0.53 \end{bmatrix} = \begin{bmatrix} 0.5925 \\ 0.3274 \end{bmatrix}$$

$$\Sigma = \begin{bmatrix} \sigma^2_{y_1} & \sigma_{y_1,y_2} \\ \sigma_{y_1,y_2} & \sigma^2_{y_2} \end{bmatrix} = \frac{1}{4} \cdot \begin{bmatrix} \sum_{i=1}^4 (y_{1i} - \mu_{y_1})^2 & \sum_{i=1}^4 (y_{1i} - \mu_{y_1})(y_{2i} - \mu_{y_2}) \\ \sum_{i=1}^4 (y_{1i} - \mu_{y_1})(y_{2i} - \mu_{y_2}) & \sum_{i=1}^4 (y_{2i} - \mu_{y_2})^2 \end{bmatrix} = \begin{bmatrix} 0.0171 & -0.0073 \\ -0.0073 & 0.0236 \end{bmatrix}$$

Now we need to compute both $|\Sigma|$ and Σ^{-1}:

$$|\Sigma| = \det \Sigma = 0.0075 \cdot 0.0075 - (-0.0025)^2 = 3.519 \cdot 10^{-4}$$

$$\Sigma^{-1} = \begin{bmatrix} 67.1101 & 20.7954 \\ 20.7954 & 48.7831 \end{bmatrix}$$

Therefore, we have the normal distribution of $\{y_1, y_2\}$ given $y_6 = B$:

$$P((y_1, y_2)|y_6 = B) = \frac{1}{\sqrt{(2\pi)^2|\Sigma|}} \exp\left(-\frac{1}{2}\left(\begin{bmatrix} y_1 \\ y_2 \end{bmatrix} - \begin{bmatrix} 0.5925 \\ 0.3274 \end{bmatrix}\right)^T \cdot \begin{bmatrix} 67.1101 & 20.7954 \\ 20.7954 & 48.7831 \end{bmatrix} \cdot \left(\begin{bmatrix} y_1 \\ y_2 \end{bmatrix} - \begin{bmatrix} 0.5925 \\ 0.3274 \end{bmatrix}\right)\right)$$

$$= \frac{1}{\sqrt{(2\pi)^2 \cdot 3.519 \cdot 10^{-4}}} \exp\left(-\frac{1}{2} \begin{bmatrix} y_1 - 0.5925 \\ y_2 - 0.3274 \end{bmatrix}^T \cdot \begin{bmatrix} 67.1101 & 20.7954 \\ 20.7954 & 48.7831 \end{bmatrix} \cdot \begin{bmatrix} y_1 - 0.5925 \\ y_2 - 0.3274 \end{bmatrix}\right)$$

Distribution of y_3 and y_4

The class conditional distributions of y_3 and y_4 come directly from the information in table **??** and they are given by:

$P(y_3 \cap y_4 | y_6 = A)$		y_3	
		0	1
y_4	0	$P(y_3 = 0 \cap y_4 = 0|y_6 = A) = 0$	$P(y_3 = 1 \cap y_4 = 0|y_6 = A) = \frac{1}{3}$
	1	$P(y_3 = 0 \cap y_4 = 1|y_6 = A) = \frac{1}{3}$	$P(y_3 = 1 \cap y_4 = 1|y_6 = A) = \frac{1}{3}$

Table 2: Distribution of y_3 and y_4 given $y_6 = A$

$P(y_3 \cap y_4 \vert y_6 = B)$		y_3	
		0	1
y_4	0	$P(y_3 = 0 \cap y_4 = 0 \vert y_6 = B) = \frac{1}{2}$	$P(y_3 = 1 \cap y_4 = 0 \vert y_6 = B) = \frac{1}{4}$
	1	$P(y_3 = 0 \cap y_4 = 1 \vert y_6 = B) = \frac{1}{4}$	$P(y_3 = 1 \cap y_4 = 1 \vert y_6 = B) = 0$

Table 3: Distribution of y_3 and y_4 given $y_6 = B$

Distribution of y_5

The class conditional distribution of y_5 is given by:

$P(y_5 \vert y_6)$		y_5		
		0	1	2
y_6	A	$P(y_5 = 0 \vert y_6 = A) = \frac{1}{3}$	$P(y_5 = 1 \vert y_6 = A) = \frac{1}{3}$	$P(y_5 = 2 \vert y_6 = A) = \frac{1}{3}$
	B	$P(y_5 = 0 \vert y_6 = B) = \frac{1}{4}$	$P(y_5 = 1 \vert y_6 = B) = \frac{1}{2}$	$P(y_5 = 2 \vert y_6 = B) = \frac{1}{4}$

Table 4: Distribution of y_5 given y_6

b)

In order to classify the testing observations, we will need to compute the posterior probabilities. Under a MAP assumption, the predicted class for each testing observation is the one that maximizes the posterior probability. Since we are only interested in the maximum value over all classes, we can ignore the denominator of the posterior probability formula. We have two testing observations, x_8 and x_9, and we will compute the posterior probabilities for each of them:

Posterior probabilities for x_8

This training observation has the following values for the input variables: $y_1 = 0.38$, $y_2 = 0.52$, $y_3 = 0$, $y_4 = 1$ and $y_5 = 0$.

$$P(y_6 = A \vert x_8) = \frac{P(x_8 \vert y_6 = A) \cdot P(y_6 = A)}{P(x_8)} \propto P(x_8 \vert y_6 = A) \cdot P(y_6 = A) =$$

$$= P(y_1 = 0.38, y_2 = 0.52 \vert y_6 = A) \cdot P(y_3 = 0, y_4 = 1 \vert y_6 = A) \cdot P(y_5 = 0 \vert y_6 = A) \cdot P(y_6 = A) =$$

$$= \frac{1}{\sqrt{(2\pi)^2 \cdot 5.4613 \cdot 10^{-5}}} \exp\left(-\frac{1}{2} \begin{bmatrix} 0.38 - 0.24 \\ 0.52 - 0.52 \end{bmatrix}^T \cdot \begin{bmatrix} 410.156 & -117.188 \\ -117.188 & 78.125 \end{bmatrix} \cdot \begin{bmatrix} 0.38 - 0.24 \\ 0.52 - 0.52 \end{bmatrix}\right) \cdot \frac{1}{3} \cdot \frac{1}{3} \cdot \frac{3}{7} =$$

$$= 0.3868 \cdot \frac{1}{3} \cdot \frac{1}{3} \cdot \frac{3}{7} = 0.01842$$

$$P(y_6 = B|x_8) = \frac{P(x_8|y_6 = B) \cdot P(y_6 = B)}{P(x_8)} \propto P(x_8|y_6 = B) \cdot P(y_6 = B) =$$

$$= P(y_1 = 0.38, y_2 = 0.52|y_6 = B) \cdot P(y_3 = 0, y_4 = 1|y_6 = B) \cdot P(y_5 = 0|y_6 = B) \cdot P(y_6 = B) =$$

$$= \frac{1}{\sqrt{(2\pi)^2 \cdot 3.519 \cdot 10^{-4}}} \exp\left(-\frac{1}{2} \begin{bmatrix} 0.38 - 0.5925 \\ 0.52 - 0.3274 \end{bmatrix}^T \cdot \begin{bmatrix} 67.1101 & 20.7954 \\ 20.7954 & 48.7831 \end{bmatrix} \cdot \begin{bmatrix} 0.38 - 0.5925 \\ 0.52 - 0.3274 \end{bmatrix}\right) \cdot \frac{1}{4} \cdot \frac{1}{4} \cdot \frac{4}{7} =$$

$$= 1.7677 \cdot \frac{1}{4} \cdot \frac{1}{4} \cdot \frac{4}{7} = 0.06313$$

Posterior probabilities for x_9

This training observation has the following values for the input variables: $y_1 = 0.42$, $y_2 = 0.59$, $y_3 = 0$, $y_4 = 1$ and $y_5 = 1$.

$$P(y_6 = A|x_9) = \frac{P(x_9|y_6 = A) \cdot P(y_6 = A)}{P(x_9)} \propto P(x_9|y_6 = A) \cdot P(y_6 = A) =$$

$$= P(y_1 = 0.42, y_2 = 0.59|y_6 = A) \cdot P(y_3 = 0, y_4 = 1|y_6 = A) \cdot P(y_5 = 1|y_6 = A) \cdot P(y_6 = A) =$$

$$= \frac{1}{\sqrt{(2\pi)^2 \cdot 5.4613 \cdot 10^{-5}}} \exp\left(-\frac{1}{2} \begin{bmatrix} 0.42 - 0.24 \\ 0.59 - 0.52 \end{bmatrix}^T \cdot \begin{bmatrix} 410.156 & -117.188 \\ -117.188 & 78.125 \end{bmatrix} \cdot \begin{bmatrix} 0.42 - 0.24 \\ 0.59 - 0.52 \end{bmatrix}\right) \cdot \frac{1}{3} \cdot \frac{1}{3} \cdot \frac{3}{7} =$$

$$= 0.1013 \cdot \frac{1}{3} \cdot \frac{1}{3} \cdot \frac{3}{7} = 0.00482$$

$$P(y_6 = B|x_9) = \frac{P(x_9|y_6 = B) \cdot P(y_6 = B)}{P(x_9)} \propto P(x_9|y_6 = B) \cdot P(y_6 = B) =$$

$$= P(y_1 = 0.42, y_2 = 0.59|y_6 = B) \cdot P(y_3 = 0, y_4 = 1|y_6 = B) \cdot P(y_5 = 1|y_6 = B) \cdot P(y_6 = B) =$$

$$= \frac{1}{\sqrt{(2\pi)^2 \cdot 3.519 \cdot 10^{-4}}} \exp\left(-\frac{1}{2} \begin{bmatrix} 0.42 - 0.5925 \\ 0.59 - 0.3274 \end{bmatrix}^T \cdot \begin{bmatrix} 67.1101 & 20.7954 \\ 20.7954 & 48.7831 \end{bmatrix} \cdot \begin{bmatrix} 0.42 - 0.5925 \\ 0.59 - 0.3274 \end{bmatrix}\right) \cdot \frac{1}{4} \cdot \frac{1}{2} \cdot \frac{4}{7} =$$

$$= 1.4927 \cdot \frac{1}{4} \cdot \frac{1}{2} \cdot \frac{4}{7} = 0.10662$$

Predicted classes

Organizing the posterior probabilities in a table, we have:

Observation	$P(y_6 = A|x_i)$	$P(y_6 = B|x_i)$
x_8	0.01842	**0.06313**
x_9	0.00482	**0.10662**

Table 5: Posterior probabilities for the testing observations

Therefore, the predicted class for both x_8 and x_9 is $y_6 = B$.

c)

Let's consider the following classifier with a unknown threshold θ whose value we aim to find:

$$f(x|\theta) = \begin{cases} A & \text{if } P(y_6 = A|x) > \theta \\ B & \text{if } P(y_6 = A|x) \leq \theta \end{cases}$$

Finding $P(y_6 = A|x)$

We are now working under an ML assumption, so, in order to classify each test observation we will only need the conditional distribution $P(x|y_6 = A)$ because every class has the same prior probability:

$$P(y_6 = A) = P(y_6 = B) = \frac{1}{2}$$

$$P(y_6 = A|x_8) = \frac{P(x_8|y_6 = A) \cdot P(y_6 = A)}{P(x_8)} = \frac{P(x_8|y_6 = A) \cdot \cancel{P(y_6 = A)}}{P(x_8|y_6 = A) \cdot \cancel{P(y_6 = A)} + P(x_8|y_6 = B) \cdot \cancel{P(y_6 = B)}} =$$

$$= \frac{P(x_8|y_6 = A)}{P(x_8|y_6 = A) + P(x_8|y_6 = B)} = \frac{0.04298}{0.04298 + 0.110478} = 0.280077$$

$$P(y_6 = A|x_9) = \frac{P(x_9|y_6 = A) \cdot P(y_6 = A)}{P(x_9)} = \frac{P(x_9|y_6 = A) \cdot \cancel{P(y_6 = A)}}{P(x_9|y_6 = A) \cdot \cancel{P(y_6 = A)} + P(x_9|y_6 = B) \cdot \cancel{P(y_6 = B)}} =$$

$$= \frac{P(x_9|y_6 = A)}{P(x_9|y_6 = A) + P(x_9|y_6 = B)} = \frac{0.011247}{0.011247 + 0.186588} = 0.056850$$

Accuracy   The accuracy of a classifier is given by:

$$\text{Accuracy} = \frac{TP + TN}{TP + TN + FP + FN}$$

where TP is the number of true positives, TN is the number of true negatives, FP is the number of false positives and FN is the number of false negatives. In our case, accuracy can only have three possible values: 0, 0.5 or 1: 0 when both x_8 and x_9 are misclassified, 0.5 when only one of them is misclassified and 1 when both are correctly classified. In order to maximize the accuracy, we want it to be 1. We know from table **??** that x_8 belongs to class A and x_9 belongs to class B. Therefore, in order to maximize the accuracy, our classifier needs to classify x_8 as A and x_9 as B. $f(x_8|\theta) = A$ therefore $\theta < P(y_6 = A|x_8) = 0.280077$ and $f(x_9|\theta) = B$ therefore $\theta > P(y_6 = A|x_9) = 0.056850$.

$$\theta \in [0.056850, 0.28007[$$

Any value of θ in this interval will maximize the accuracy of our classifier.

2nd Question

a)

In order to obtain y_2 under an equal-range discretization, we followed the rule:

$$y_{2_{normalized}} = \begin{cases} 0 & \text{if } y_2 \in [0, 0.5) \\ 1 & \text{if } y_2 \in [0.5, 1] \end{cases}$$

The normalized values are:

Dataset	x_1	x_2	x_3	x_4	x_5	x_6	x_7	x_8	x_9
y_2	0	0	1	0	0	0	1	1	1

Table 6: Normalized y_2 values

y_1 will now be considered an output variable and y_2 to y_6 will be considered input variables. Considering the normalized values of y_2, we can rewrite the dataset as follows:

D		Output	Input				
Fold		y_1	$y_{2_{norm}}$	y_3	y_4	y_5	y_6
F_1	x_1	0.24	0	1	1	0	A
	x_2	0.16	0	1	0	1	A
	x_3	0.32	1	0	1	2	A
F_2	x_4	0.54	0	0	0	1	B
	x_5	0.66	0	0	0	0	B
	x_6	0.76	0	1	0	2	B
F_3	x_7	0.41	1	0	1	1	B
	x_8	0.38	1	0	1	0	A
	x_9	0.42	1	0	1	1	B

Table 7: Dataset D divided into three folds

Adittionally, we have divided the dataset into three folds, F_1, F_2 and F_3.

b)

In this exercise we aim to compute a kNN (k nearest neighbors - Lazy Learning) classifier considering the following parameters:

- $k = 3$

- **Hamming Distance** as the distance to be used to compute the nearest neighbors of a given observation.

$$d_H(x_i, x_j) = \sum_{l=1}^{m} \delta(y_{il}, y_{jl})$$

where m is the number of input variables and y_{il} is the value of the l^{th} input variable of the i^{th} observation.

We have divided our dataset in folds in order to perform a cross validation. We will only be interested in the first iteration of the cross validation, where F_3 is the testing fold and F_1 and F_2 are the training folds. We will now compute the kNN classifier for each observation in F_3 and afterwards we will compute the MAE (Mean Absolute Error) for the testing fold.

Computing the Hamming Distances

Testing Observation (x_i)	$d_H(x_i, x_1)$	$d_H(x_i, x_2)$	$d_H(x_i, x_3)$	$d_H(x_i, x_4)$	$d_H(x_i, x_5)$	$d_H(x_i, x_6)$
x_7	4	4	2	2	3	4
x_8	2	4	1	4	3	5
x_9	4	4	2	2	3	4

Table 8: Hamming distances between the testing observation x_i and the training observations x_j

We are considering $k = 3$, so we will only need the three nearest neighbors of each testing observation. In the table above we have filled with yellow the three nearest neighbors of each testing observation.

Predicted value of y_1 for each testing observation

The output value we are working with is numerical, so the predicted value of y_1 for each testing observation will be:

$$\hat{y}_{1j} = \frac{\sum_{i=1}^{k} \frac{1}{d_H(x_i, x_j)} \cdot y_{1j}}{\sum_{i=1}^{k} \frac{1}{d_H(x_i, x_j)}}$$

where k is the number of nearest neighbors, $d_H(x_i, x_j)$ is the Hamming distance between the testing observation x_i and the i^{th} nearest neighbor and y_{1j} is the value of the output variable of the j^{th} nearest neighbor. $\frac{1}{d_H(x_i, x_j)}$ is the weight of the i^{th} nearest neighbor.

$$\hat{y}_{17} = \frac{1}{\frac{1}{2} + \frac{1}{2} + \frac{1}{3}} \cdot \left(\frac{1}{2} \cdot 0.32 + \frac{1}{2} \cdot 0.54 + \frac{1}{3} \cdot 0.66 \right) = 0.4875$$

$$\hat{y}_{18} = \frac{1}{\frac{1}{2} + \frac{1}{1} + \frac{1}{3}} \cdot \left(\frac{1}{2} \cdot 0.24 + 1 \cdot 0.32 + \frac{1}{3} \cdot 0.66 \right) = 0.36$$

$$\hat{y}_{19} = \frac{1}{\frac{1}{2} + \frac{1}{2} + \frac{1}{3}} \cdot \left(\frac{1}{2} \cdot 0.32 + \frac{1}{2} \cdot 0.54 + \frac{1}{3} \cdot 0.66 \right) = 0.4875$$

We have the following predicted values for y_1:

Testing Observation (x_i)	$\hat{y}_{1i}$	y_{1i}
x_7	0.4875	0.41
x_8	0.36	0.38
x_9	0.4875	0.42

Table 9: Predicted values of y_1 for each testing observation

MAE

The MAE is given by:

$$\text{MAE} = \frac{1}{n} \sum_{i=1}^{n} |y_{1i} - \hat{y}_{1i}|$$

where n is the number of testing observations, y_{1i} is the value of the output variable of the i^{th} testing observation and $\hat{y}_{1i}$ is the predicted value of the output variable of the i^{th} testing observation.

$$\text{MAE} = \frac{1}{3} \cdot \left(|0.41 - 0.4875| + |0.38 - 0.36| + |0.42 - 0.4875| \right) = 0.055$$

Programming and Critical Analysis

Imports

```
1   # Sklearn Imports
2   import sklearn as sk
3   from sklearn.metrics import confusion_matrix
4   from sklearn.model_selection import cross_val_score, StratifiedKFold
5   from sklearn.neighbors import KNeighborsClassifier
6   from sklearn.naive_bayes import GaussianNB
7   # Other Imports
8   import numpy as np
9   import pandas as pd
10  import matplotlib.pyplot as plt
11  import seaborn as sns
12  from pathlib import Path
13  from scipy.io.arff import loadarff
14  from scipy import stats
```

Listing 1:

Loading DataSet and Define Directories

```
1   IMAGES_DIR = Path('images')
2   IMAGES_DIR.mkdir(parents=True, exist_ok=True)
3   DATA_DIR = Path('data')
4   DATA_DIR.mkdir(parents=True, exist_ok=True)
5   DATA_FILE = 'column_diagnosis.arff'
6   DATA_PATH = DATA_DIR / DATA_FILE
7   data = loadarff(DATA_PATH)
8   df = pd.DataFrame(data[0])
9   df['class'] = df['class'].str.decode('utf-8')
10  # Show the first 5 rows
11  df.head()
```

Listing 2:

1ˢᵗ Question

10-fold cross validation with suffling

```
1   # Split into features and labels
2   X = df.drop('class', axis=1)
3   y = df['class']
4
```

Listing 3:

```
1   # Stratified 10 fold cross validation
2   cv = StratifiedKFold(n_splits=10, shuffle=True, random_state=0)
3
```

Listing 4:

Box Plots for Fold Accuracy

```
1    # kNN Classifier with 5 neighbors
2    knn = KNeighborsClassifier(n_neighbors=5)
3
4    # Naive Bayes Classifier
5    nb = GaussianNB()
6
7    # 10-fold startified cross validation with shuffle
8    knn_score = cross_val_score(knn,
9                                X,
10                               y,
11                               cv=cv,
12                               scoring='accuracy')
13
14   nb_score = cross_val_score(nb,
15                              X,
16                              y,
17                              cv=cv,
18                              scoring='accuracy')
19
```

Listing 5:

```
1    # Create boxplots for the accuracies of both classifiers
2    plt.boxplot([knn_score, nb_score], labels=['k-NN', 'Naive Bayes'])
3    plt.ylabel('Accuracy')
4    plt.title('Accuracy Comparison Between k-NN and Naive Bayes')
5    plt.savefig(IMAGES_DIR / 'boxplot.png')
6    plt.show()
7
```

Listing 6:
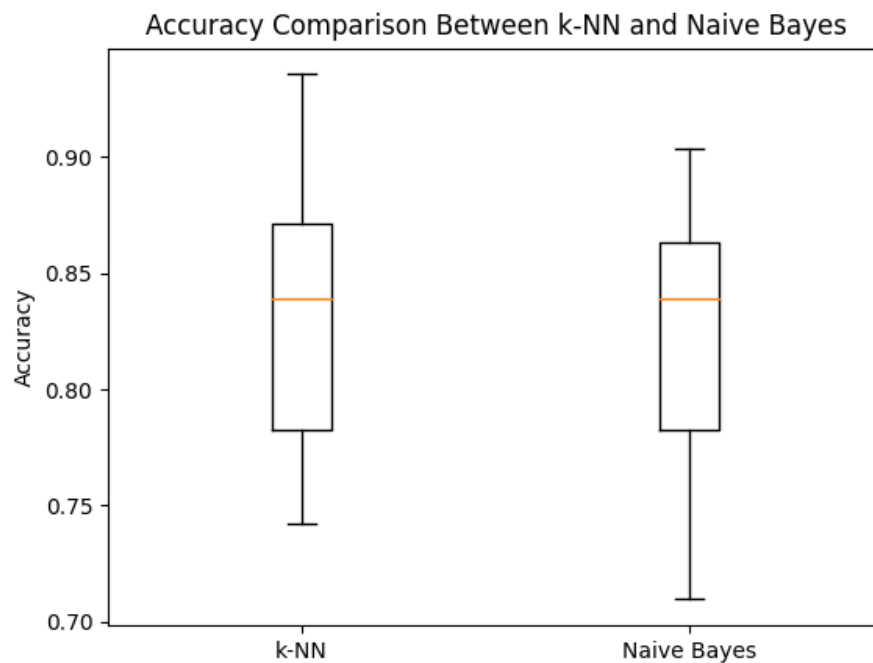
Figure 1: Boxplot for the accuracies of both classifiers

Comment: As seen by the box plot above, both models have an equal central tendency, shown by the median. Beyond that, the variance of the models is quite similar, since the overall amplitude of the box plots is similar. However, Naive Bayes has a variance towards the lower end of the box plot, meaning it has a tendency for lower accuracy.

```python
    # Perform a two-sample t-test to compare the means
t_stat, p_value = stats.ttest_rel(knn_score, nb_score, alternative='greater')
print(f"t-statistic: {str(t_stat)}")
print(f"p-value: {str(p_value)}")

# Determine if the null hypothesis can be rejected (p < 0.05)
alpha = 0.05
if p_value < alpha:
    print("Null hypothesis rejected: \nkNN is statistically superior to Naive Bayes.")
else:
    print("Null hypothesis not rejected: \nNo significant difference between kNN and
    Naive Bayes.")
```

Listing 7:

```
    t-statistic: 0.9214426752509264
    p-value: 0.19042809062064092
    Null hypothesis not rejected:
    No significant difference between kNN and Naive Bayes.
```

Listing 8:

2ⁿᵈ Question

```python
# Create k-NN classifiers with k=1 and k=5
knn_k1 = KNeighborsClassifier(n_neighbors=1, weights='uniform', metric='euclidean')
knn_k5 = KNeighborsClassifier(n_neighbors=5, weights='uniform', metric='euclidean')

# Create a confusion matrix for each classifier
confusion_matrix_k1 = np.zeros((3, 3))
confusion_matrix_k5 = np.zeros((3, 3))

for train_index, test_index in cv.split(X, y):
    # Split the data into training and testing sets
    X_train, X_test = X.iloc[train_index], X.iloc[test_index]
    y_train, y_test = y.iloc[train_index], y.iloc[test_index]

    # Fit both models to the dataset
    knn_k1.fit(X_train, y_train)
    knn_k5.fit(X_train, y_train)

    # Make predictions using both models
    predictions_k1 = knn_k1.predict(X_test)
    predictions_k5 = knn_k5.predict(X_test)

    # Add the confusion matrices to the cumulative confusion matrices
    confusion_matrix_k1 += confusion_matrix(y_test, predictions_k1)
    confusion_matrix_k5 += confusion_matrix(y_test, predictions_k5)

# Plot the differences between the cumulative confusion matrices
# Get the class names
class_names = np.unique(y)

# Calculate the difference between the cumulative confusion matrices
confusion_matrix_difference = confusion_matrix_k1 - confusion_matrix_k5
```

```
32
33   # Plot the differences between the cumulative confusion matrices
34   plt.figure(figsize=(8, 6))
35   plt.imshow(confusion_matrix_difference, cmap='coolwarm', interpolation='nearest')
36   plt.colorbar(label='Difference')
37   plt.title('Confusion Matrix Difference (k=1 - k=5)')
38   plt.xticks(ticks=np.arange(len(class_names)), labels=class_names)
39   plt.yticks(ticks=np.arange(len(class_names)), labels=class_names)
40   plt.xlabel('Predicted Class')
41   plt.ylabel('True Class')
42
43   # Loop over the cells and add the values to the plot
44   for i in range(len(class_names)):
45       for j in range(len(class_names)):
46           plt.text(j, i, str(confusion_matrix_difference[i, j]), horizontalalignment='
     center', verticalalignment='center')
47
48   plt.savefig(IMAGES_DIR / 'confusion_matrix_difference.png')
49   plt.show()
```

Listing 9:

Figure 2: Confusion Matrix Difference (k=1 - k=5)

Comment: As seen by the difference confusion cummulative matrix above, the classifiers do not have significant differences, since the highest difference between them is 5. As for some entraces of the matrix, one can say both classifiers have the same performance, since the difference between them is 0. For negative values on the diagonal, it means that the classifier knn5 has a better performance than the classifier knn1, and for positive values, the opposite. As one can see, what stands out prominently in the matrix is the prevalence of positive values, outnumbering the negative ones. This observation implies that k-NN5 demonstrates superior performance in distinguishing between classes and encounters fewer instances of confusion compared to k-NN1.

3rd Question

There are some problems that Naive Bayes may encounter. The tree main problems we can think of are:

- Independence Assumption

- Continuous Features

- Imbalanced Classes

Independence Assumption

Naive Bayes assumes that features are conditionally independent given the class label. In other words, it assumes that there are no correlations between the features. However, in real-world datasets like medical diagnoses, features may be correlated. For example, certain symptoms or medical test results might be related. If there are strong correlations between features, Naive Bayes may not capture these relationships accurately, leading to suboptimal performance.

Continuous Feature

If the dataset contains continuous or numerical features, Naive Bayes relies on assuming that the data follows a specific probability distribution (e.g., Gaussian for Gaussian Naive Bayes). If the data distribution significantly deviates from this assumption, the model's performance can degrade. Medical datasets often contain a mix of continuous and categorical features, and handling continuous data can be a challenge for Naive Bayes.

Imbalanced Classes

If the dataset has imbalanced class distributions, where one class significantly outnumbers the others, Naive Bayes can be biased towards the majority class. This is because the class prior probabilities heavily influence the classification decision. In medical datasets, it's common to have imbalanced class distributions, such as a rare disease. Naive Bayes may struggle to correctly classify the minority class due to the bias.

To address if those problems exist in our dataset, we can do the following:

- Check if there are correlations between features

- Check if the data follows the assumed probability distribution

- Check if the class distributions are balanced

We do that in the following code:

Independence Assumption

```
# Independence Assumption
# Calculate pairwise feature correlations
correlation_matrix = X.corr()

# Visualize feature correlations using a heatmap
plt.figure(figsize=(10, 8))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', fmt=".2f")
plt.title("Feature Correlation Heatmap")
plt.savefig(IMAGES_DIR / 'correlation_heatmap.png')
plt.show()
```

Listing 10:

Figure 3: Feature Correlation Heatmap

Comment: As one can see in the correlation matrix, there are some features that are correlated. For example, the features "sacral slope" and "pelvic incidence" are correlated with a correlation coefficient of 0.81, which is considered as a strong correlation. This means that Naive Bayes may not be the best choice for this dataset, since it assumes that features are independent.

Continuous Features

```python
# Continuous Features
# Identify continuous features (assuming all non-integer features are continuous)
continuous_features = [col for col in X.columns if X[col].dtype != 'int64']

# Plot histograms for continuous features
num_cols = 3
num_rows = int(np.ceil(len(continuous_features) / num_cols))
fig, axs = plt.subplots(num_rows, num_cols, figsize=(20, 4*num_rows))

for i, feature in enumerate(continuous_features):
    row = i // num_cols
    col = i % num_cols
    sns.histplot(data=X, x=feature, kde=True, color='skyblue', ax=axs[row][col])
    axs[row][col].set_title(f'Distribution of {feature}')

plt.savefig(IMAGES_DIR / 'continuous_features.png')
plt.show()
```
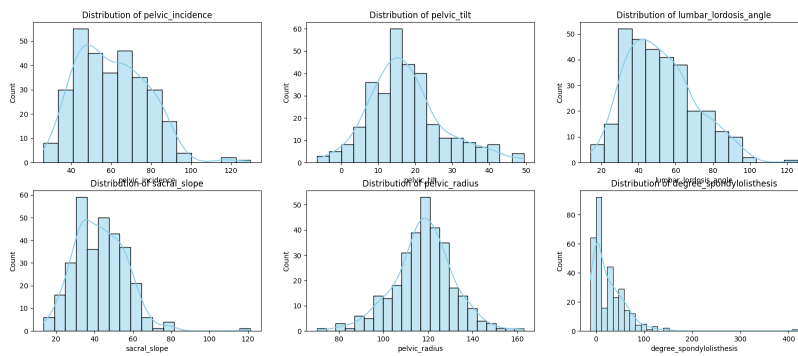
Listing 11:

Figure 4: Distribution of continuous features

Comment: As one can see in the above plot, almost all features resemble a Gaussian distribution. However, there are some features that do not follow a Gaussian distribution. For example, the Degree Spondylolisthesis is close to the origin, more like a Gamma distribution. This may cause problems for Naive Bayes, since it assumes a Gaussian distribution. To resolve this possible issue, one can try to preprocess the data in order to make it more Gaussian-like.

Imbalanced Classes

```python
# Imbalanced Classes
# Count class frequencies
class_counts = y.value_counts()

# Visualize class distribution
plt.figure(figsize=(6, 4))
sns.barplot(x=class_counts.index, y=class_counts.values, palette="Set2")
plt.title("Class Distribution")
plt.xlabel("Class")
plt.ylabel("Count")
plt.savefig(IMAGES_DIR / 'class_distribution.png')
plt.show()

# Output class frequencies
print("Class Frequencies:")
print(class_counts)
```

Listing 12:

Figure 5: Class Distribution

```
1   Class Frequencies:
2   class
3   Spondylolisthesis    150
4   Normal               100
5   Hernia                60
6   Name: count, dtype: int64
```

Listing 13:

Comment: As one can see in the above plot, the class distributions do not seem to be balanced. The Spondylolisthesis class significantly outnumbers the other two classes, making it more than the double of the Hernia class and 50% more than the Normal Class. This may cause problems for Naive Bayes, since it can be biased towards the majority class. To resolve this possible issue, one can try to balance the class distributions by either undersampling the majority class or oversampling the minority classes.